

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

In [155]: df=pd.read_csv(r"C:\Users\Admin\Downloads\csvs_per_year\csvs_per_year\madrid_2005.csv")
df
```

Out[155]:

	date	BEN	CO	EBE	MXV	NMHC	NO_2	NOx	OXY	O_3	PM10	PM25	PXY	SO_2
0	2005-11-01 01:00:00	NaN	0.77	NaN	NaN	NaN	57.130001	128.699997	NaN	14.720000	14.91	10.65	NaN	4.6
1	2005-11-01 01:00:00	1.52	0.65	1.49	4.57	0.25	86.559998	181.699997	1.27	11.680000	30.93	NaN	1.59	7.8
2	2005-11-01 01:00:00	NaN	0.40	NaN	NaN	NaN	46.119999	53.000000	NaN	30.469999	14.60	NaN	NaN	5.7
3	2005-11-01 01:00:00	NaN	0.42	NaN	NaN	NaN	37.220001	52.009998	NaN	21.379999	15.16	NaN	NaN	6.6
4	2005-11-01 01:00:00	NaN	0.57	NaN	NaN	NaN	32.160000	36.680000	NaN	33.410000	5.00	NaN	NaN	3.0
...	...	...	...	...	...	...	...	...	...	...	...	...	...	.
236995	2006-01-01 00:00:00	1.08	0.36	1.01	NaN	0.11	21.990000	23.610001	NaN	43.349998	5.00	NaN	NaN	6.6
236996	2006-01-01 00:00:00	0.39	0.54	1.00	1.00	0.11	2.200000	4.220000	1.00	69.639999	4.95	1.49	1.00	7.0
236997	2006-01-01 00:00:00	0.19	NaN	0.26	NaN	0.08	26.730000	30.809999	NaN	43.840000	4.31	2.93	NaN	13.2
236998	2006-01-01 00:00:00	0.14	NaN	1.00	NaN	0.06	13.770000	17.770000	NaN	NaN	5.00	NaN	NaN	5.8
236999	2006-01-01 00:00:00	0.50	0.40	0.73	1.84	0.13	20.940001	26.950001	1.49	48.259998	5.67	2.11	1.09	11.0

237000 rows × 17 columns

```
In [156]: df1 = df.fillna(0)
df1
```

Out[156]:

	date	BEN	CO	EBE	MXV	NMHC	NO_2	NOx	OXY	O_3	PM10	PM25	PXY	SO_2
0	2005-11-01 01:00:00	0.00	0.77	0.00	0.00	0.00	57.130001	128.699997	0.00	14.720000	14.91	10.65	0.00	4.61
1	2005-11-01 01:00:00	1.52	0.65	1.49	4.57	0.25	86.559998	181.699997	1.27	11.680000	30.93	0.00	1.59	7.81
2	2005-11-01 01:00:00	0.00	0.40	0.00	0.00	0.00	46.119999	53.000000	0.00	30.469999	14.60	0.00	0.00	5.71
3	2005-11-01 01:00:00	0.00	0.42	0.00	0.00	0.00	37.220001	52.009998	0.00	21.379999	15.16	0.00	0.00	6.61
4	2005-11-01 01:00:00	0.00	0.57	0.00	0.00	0.00	32.160000	36.680000	0.00	33.410000	5.00	0.00	0.00	3.01
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
236995	2006-01-01 00:00:00	1.08	0.36	1.01	0.00	0.11	21.990000	23.610001	0.00	43.349998	5.00	0.00	0.00	6.61
236996	2006-01-01 00:00:00	0.39	0.54	1.00	1.00	0.11	2.200000	4.220000	1.00	69.639999	4.95	1.49	1.00	7.01
236997	2006-01-01 00:00:00	0.19	0.00	0.26	0.00	0.08	26.730000	30.809999	0.00	43.840000	4.31	2.93	0.00	13.21
236998	2006-01-01 00:00:00	0.14	0.00	1.00	0.00	0.06	13.770000	17.770000	0.00	0.000000	5.00	0.00	0.00	5.81
236999	2006-01-01 00:00:00	0.50	0.40	0.73	1.84	0.13	20.940001	26.950001	1.49	48.259998	5.67	2.11	1.09	11.01

237000 rows × 17 columns

In [157]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 237000 entries, 0 to 236999
Data columns (total 17 columns):
#   Column      Non-Null Count  Dtype
---  -
0   date        237000 non-null  object
1   BEN         70370 non-null   float64
2   CO          217656 non-null   float64
3   EBE         68955 non-null   float64
4   MXY         32549 non-null   float64
5   NMHC        92854 non-null   float64
6   NO_2        235022 non-null   float64
7   NOx         235049 non-null   float64
8   OXY         32555 non-null   float64
9   O_3         223162 non-null   float64
10  PM10        232142 non-null   float64
11  PM25        69407 non-null   float64
12  PXY         32549 non-null   float64
13  SO_2        235277 non-null   float64
14  TCH         93076 non-null   float64
15  TOL         70255 non-null   float64
16  station     237000 non-null   int64
dtypes: float64(15), int64(1), object(1)
memory usage: 30.7+ MB
```

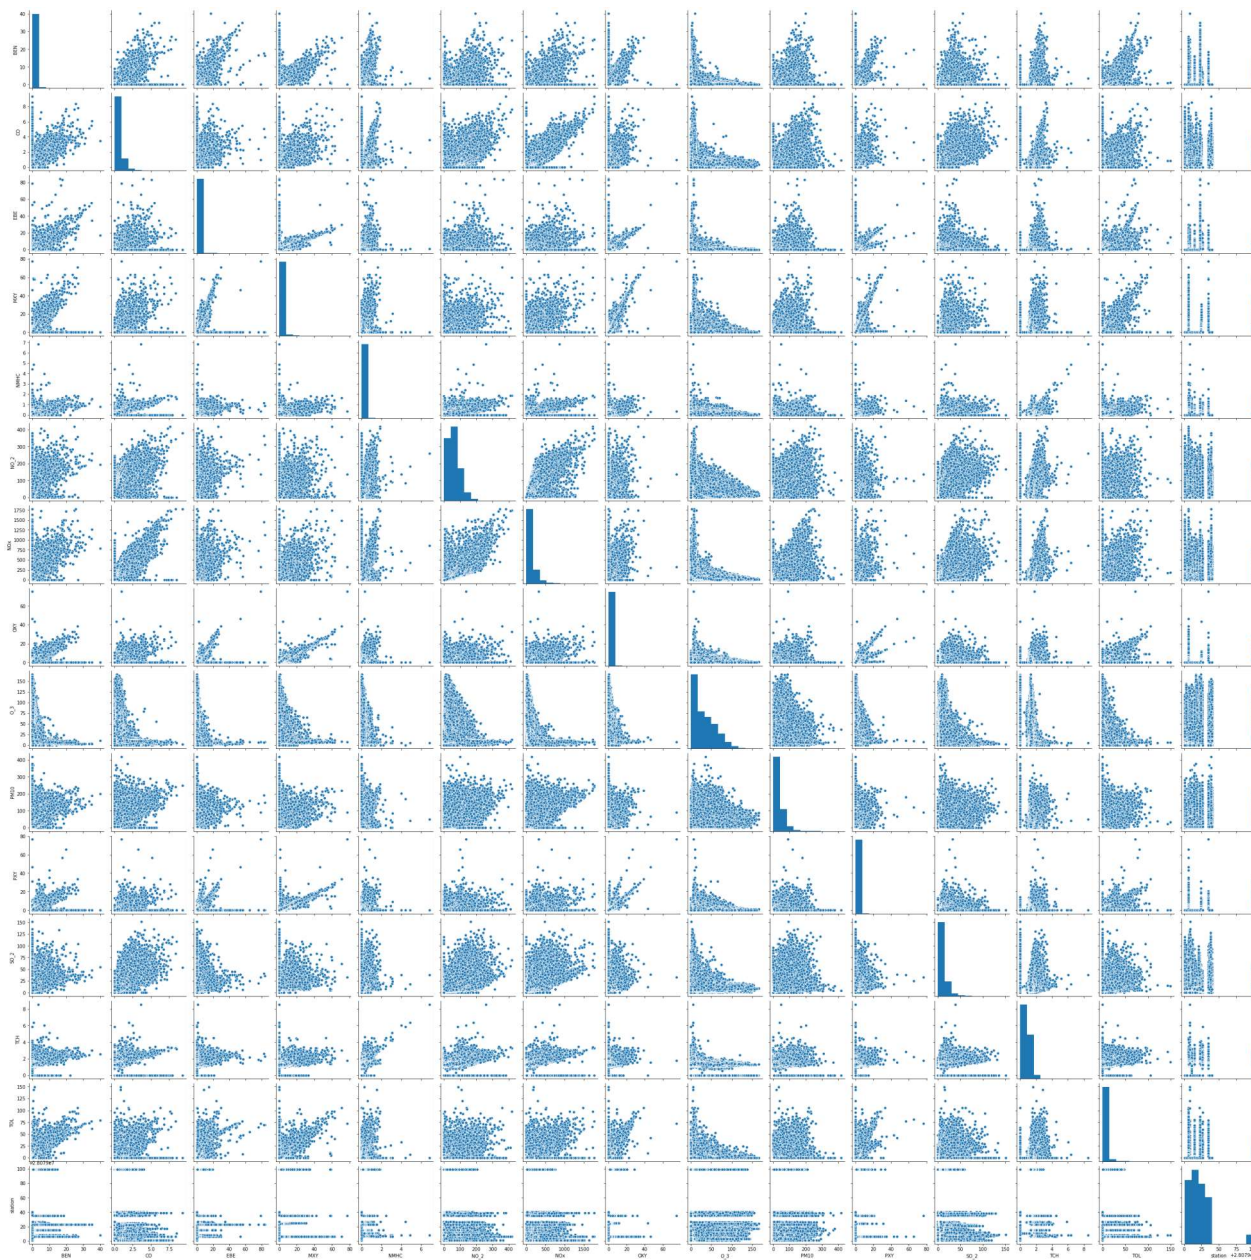
In [158]: `df.columns`

Out[158]: Index(['date', 'BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO\_2', 'NOx', 'OXY', 'O\_3', 'PM10', 'PM25', 'PXY', 'SO\_2', 'TCH', 'TOL', 'station'], dtype='object')

In [159]: `df2=df1[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3', 'PM10', 'PXY', 'SO_2', 'TCH', 'TOL', 'station']]`


```
In [160]: sns.pairplot(df2)
```

```
Out[160]: <seaborn.axisgrid.PairGrid at 0x23344056f70>
```




```
In [161]: sns.distplot(df2['O_3'])
```

```
Out[161]: <matplotlib.axes._subplots.AxesSubplot at 0x2343e148fa0>
```



```
In [162]: sns.heatmap(df2.corr())
```

```
Out[162]: <matplotlib.axes._subplots.AxesSubplot at 0x234566f54c0>
```



## Linear Regression

```
In [163]: x = df2[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY',
                  'PM10', 'PXY', 'SO_2', 'TCH', 'TOL']]
y = df2['O_3']
```

```
In [164]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.3)
```

```
In [165]: from sklearn.linear_model import LinearRegression

lr=LinearRegression()
lr.fit(x_train,y_train)
```

```
Out[165]: LinearRegression()
```

```
In [166]: print(lr.intercept_)
```

```
51.42294040923451
```


```
In [167]: coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-effecient'])
coeff
```

```
Out[167]:
```

	Co-effecient
<b>BEN</b>	1.552194
<b>CO</b>	4.035896
<b>EBE</b>	-0.026706
<b>MXV</b>	-0.382310
<b>NMHC</b>	7.055252
<b>NO_2</b>	-0.225619
<b>NOx</b>	-0.085837
<b>OXY</b>	0.856665
<b>PM10</b>	0.113306
<b>PXY</b>	0.538582
<b>SO_2</b>	-0.032458
<b>TCH</b>	-2.506577
<b>TOL</b>	-0.277806

```
In [168]: prediction=lr.predict(x_test)
plt.scatter(y_test,prediction)
```

```
Out[168]: <matplotlib.collections.PathCollection at 0x2344f9a4b50>
```



```
In [169]: print(lr.score(x_test,y_test))
```

```
0.28903819112398743
```

```
In [170]: lr.score(x_train,y_train)
```

```
Out[170]: 0.28783184644650417
```

## Ridge Lasso

```
In [171]: from sklearn.linear_model import Ridge,Lasso
```

```
In [172]: rr=Ridge(alpha=10)
rr.fit(x_train,y_train)
rr.score(x_test,y_test)
```

```
Out[172]: 0.28903742721761205
```

```
In [173]: predict2=(rr.predict(x_test))
```

```
In [174]: la=Lasso(alpha=10)
la.fit(x_train,y_train)
```

```
Out[174]: Lasso(alpha=10)
```

```
In [175]: la.score(x_test,y_test)
```

```
Out[175]: 0.28085040815849505
```

## Elastic Net regression

```
In [176]: from sklearn.linear_model import ElasticNet
en=ElasticNet()
en.fit(x_train,y_train)
```

```
Out[176]: ElasticNet()
```

```
In [177]: print(en.coef_)
```

```
[ 0.00774178  0.          0.          0.14495589 -0.          -0.23862087
 -0.06696927  0.          0.11308583  0.0986593   0.03145642 -0.33690942
 -0.0502186 ]
```

```
In [178]: print(en.intercept_)
```

```
51.13160182619739
```

```
In [179]: print(en.score(x_test,y_test))
```

```
0.28285699603945946
```

```
In [180]: print(en.score(x_train,y_train))
```

```
0.28216854923305
```

## Logistic Regression

```
In [181]: from sklearn.linear_model import LogisticRegression
```

```
In [182]: feature_matrix=df2.iloc[:,0:5]
target_vector=df2.iloc[:,-1]
```

```
In [183]: from sklearn.preprocessing import StandardScaler
```

```
In [184]: fs=StandardScaler().fit_transform(feature_matrix)
```

```
In [185]: logr=LogisticRegression()  
logr.fit(fs,target_vector)
```

C:\Users\Admin\anaconda3\lib\site-packages\sklearn\linear\_model\\_logistic.py:762: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression) ([https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression))

```
n_iter_i = _check_optimize_result(
```

```
Out[185]: LogisticRegression()
```

```
In [186]: df2.shape
```

```
Out[186]: (237000, 15)
```

```
In [187]: observation=[[1,2,3,4,5]]  
predication = logr.predict(observation)
```

```
In [188]: print(predication)
```

```
[28079099]
```

```
In [189]: logr.classes_
```

```
Out[189]: array([28079001, 28079003, 28079004, 28079006, 28079007, 28079008,  
                28079009, 28079011, 28079012, 28079014, 28079015, 28079016,  
                28079017, 28079018, 28079019, 28079021, 28079022, 28079023,  
                28079024, 28079025, 28079026, 28079027, 28079035, 28079036,  
                28079038, 28079039, 28079040, 28079099], dtype=int64)
```

```
In [190]: from sklearn.model_selection import train_test_split
```

```
x_train,x_test,y_train,y_test = train_test_split(feature_matrix,target_vector,test_size=0.30
```

```
In [191]: print(logr.score(x_test,y_test))
```

```
0.10585091420534458
```

```
In [192]: print(logr.score(x_train,y_train))
```

```
0.10491259795057263
```

## Conclusion

Linear Regression is bestfit model

The Score x\_test,y\_test is 0.28903819112398743 and x\_train,y\_train score is 0.28783184644650417



In [ ]:

In [ ]: