

```
In [6]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [45]: df=pd.read_csv(r"C:\Users\Admin\Downloads\csvs_per_year\csvs_per_year\madrid_2010.csv")
df
```

Out[45]:

	date	BEN	CO	EBE	MXV	NMHC	NO_2	NOx	OXY	O_3	PM10
0	2010-03-01 01:00:00	NaN	0.29	NaN	NaN	NaN	25.090000	29.219999	NaN	68.930000	19.410000
1	2010-03-01 01:00:00	NaN	0.27	NaN	NaN	NaN	24.879999	30.040001	NaN	NaN	19.410000
2	2010-03-01 01:00:00	NaN	0.28	NaN	NaN	NaN	17.410000	20.540001	NaN	72.120003	19.410000
3	2010-03-01 01:00:00	0.38	0.24	1.74	NaN	0.05	15.610000	21.080000	NaN	72.970001	19.410000
4	2010-03-01 01:00:00	0.79	NaN	1.32	NaN	NaN	21.430000	26.070000	NaN	NaN	24.670000
...
209443	2010-08-01 00:00:00	NaN	0.55	NaN	NaN	NaN	125.000000	219.899994	NaN	25.379999	19.410000
209444	2010-08-01 00:00:00	NaN	0.27	NaN	NaN	NaN	45.709999	47.410000	NaN	NaN	51.259000
209445	2010-08-01 00:00:00	NaN	NaN	NaN	NaN	0.24	46.560001	49.040001	NaN	46.250000	19.410000
209446	2010-08-01 00:00:00	NaN	NaN	NaN	NaN	NaN	46.770000	50.119999	NaN	77.709999	19.410000
209447	2010-08-01 00:00:00	0.92	0.43	0.71	NaN	0.25	76.330002	88.190002	NaN	52.259998	47.150000

209448 rows × 12 columns

```
In [46]: df1 = df.fillna(0)
df1
```

Out[46]:

	date	BEN	CO	EBE	MXV	NMHC	NO_2	NOx	OXY	O_3	PI
0	2010-03-01 01:00:00	0.00	0.29	0.00	0.0	0.00	25.090000	29.219999	0.0	68.930000	0.000
1	2010-03-01 01:00:00	0.00	0.27	0.00	0.0	0.00	24.879999	30.040001	0.0	0.000000	0.000
2	2010-03-01 01:00:00	0.00	0.28	0.00	0.0	0.00	17.410000	20.540001	0.0	72.120003	0.000
3	2010-03-01 01:00:00	0.38	0.24	1.74	0.0	0.05	15.610000	21.080000	0.0	72.970001	19.410
4	2010-03-01 01:00:00	0.79	0.00	1.32	0.0	0.00	21.430000	26.070000	0.0	0.000000	24.670
...
209443	2010-08-01 00:00:00	0.00	0.55	0.00	0.0	0.00	125.000000	219.899994	0.0	25.379999	0.000
209444	2010-08-01 00:00:00	0.00	0.27	0.00	0.0	0.00	45.709999	47.410000	0.0	0.000000	51.259
209445	2010-08-01 00:00:00	0.00	0.00	0.00	0.0	0.24	46.560001	49.040001	0.0	46.250000	0.000
209446	2010-08-01 00:00:00	0.00	0.00	0.00	0.0	0.00	46.770000	50.119999	0.0	77.709999	0.000
209447	2010-08-01 00:00:00	0.92	0.43	0.71	0.0	0.25	76.330002	88.190002	0.0	52.259998	47.150

209448 rows × 17 columns

In [47]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 209448 entries, 0 to 209447
Data columns (total 17 columns):
#   Column      Non-Null Count  Dtype
---  -
0   date        209448 non-null object
1   BEN         60268 non-null float64
2   CO          94982 non-null float64
3   EBE         60253 non-null float64
4   MXY         6750 non-null  float64
5   NMHC        51727 non-null float64
6   NO_2        208219 non-null float64
7   NOx         208210 non-null float64
8   OXY         6750 non-null  float64
9   O_3         126684 non-null float64
10  PM10        106186 non-null float64
11  PM25        55514 non-null float64
12  PXY         6740 non-null  float64
13  SO_2        93184 non-null float64
14  TCH         51730 non-null float64
15  TOL         60171 non-null float64
16  station     209448 non-null int64
dtypes: float64(15), int64(1), object(1)
memory usage: 27.2+ MB
```

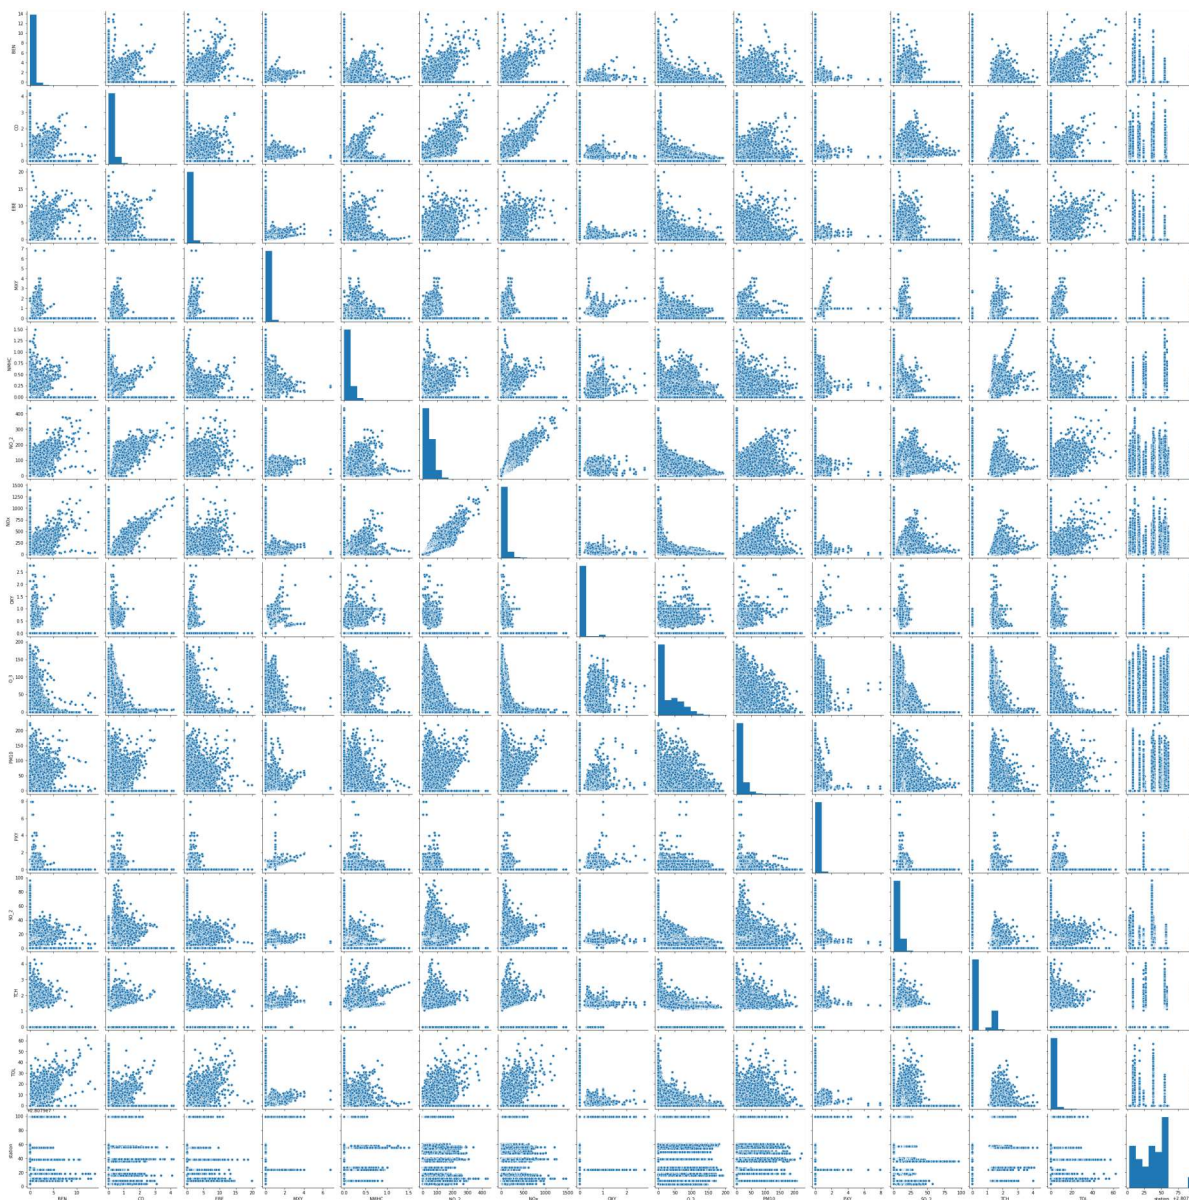
In [48]: `df.columns`

Out[48]: Index(['date', 'BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
'PM10', 'PM25', 'PXY', 'SO_2', 'TCH', 'TOL', 'station'],
dtype='object')

In [49]: `df2=df1[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
'PM10', 'PXY', 'SO_2', 'TCH', 'TOL', 'station']]`


```
In [50]: sns.pairplot(df2)
```

```
Out[50]: <seaborn.axisgrid.PairGrid at 0x1ac0db62820>
```




```
In [51]: sns.distplot(df2['O_3'])
```

```
Out[51]: <matplotlib.axes._subplots.AxesSubplot at 0x1ac5f63f1f0>
```



```
In [52]: sns.heatmap(df2.corr())
```

```
Out[52]: <matplotlib.axes._subplots.AxesSubplot at 0x1ac5ff29670>
```



Linear Regression

```
In [53]: x = df2[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY',
                'PM10', 'PXY', 'SO_2', 'TCH', 'TOL']]
y = df2['O_3']
```

```
In [54]: from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.3)
```

```
In [55]: from sklearn.linear_model import LinearRegression
```

```
lr=LinearRegression()  
lr.fit(x_train,y_train)
```

```
Out[55]: LinearRegression()
```

```
In [56]: print(lr.intercept_)
```

```
47.43549509386216
```


```
In [57]: coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-effecient'])  
coeff
```

```
Out[57]:
```

	Co-effecient
BEN	-3.270166
CO	27.242984
EBE	3.883891
MXY	-3.514638
NMHC	-48.684393
NO_2	-0.404697
NOx	-0.016305
OXY	30.458939
PM10	-0.245842
PXY	-4.226543
SO_2	-1.182267
TCH	14.136102
TOL	0.015777

```
In [58]: prediction=lr.predict(x_test)  
plt.scatter(y_test,prediction)
```

```
Out[58]: <matplotlib.collections.PathCollection at 0x1ac61185c70>
```



```
In [59]: print(lr.score(x_test,y_test))
```

```
0.258645747289049
```

```
In [60]: lr.score(x_train,y_train)
```

```
Out[60]: 0.2571856671711411
```

Ridge Lasso

```
In [61]: from sklearn.linear_model import Ridge,Lasso
```

```
In [62]: rr=Ridge(alpha=10)
rr.fit(x_train,y_train)
rr.score(x_test,y_test)
```

```
Out[62]: 0.25865307425094985
```

```
In [63]: predict2=(rr.predict(x_test))
```

```
In [64]: la=Lasso(alpha=10)
la.fit(x_train,y_train)
```

```
Out[64]: Lasso(alpha=10)
```

```
In [65]: la.score(x_test,y_test)
```

```
Out[65]: 0.18152281962192107
```

Elastic Net regression

```
In [66]: from sklearn.linear_model import ElasticNet
en=ElasticNet()
en.fit(x_train,y_train)
```

```
Out[66]: ElasticNet()
```

```
In [67]: print(en.coef_)
```

```
[ 0.          0.8909809  1.49547789  0.19498996  0.         -0.43545558
  0.02069334  0.38943572 -0.25780997  0.14675393 -0.52754491  3.376013
  0.          ]
```

```
In [68]: print(en.intercept_)
```

```
49.839760307248845
```

```
In [69]: print(en.score(x_test,y_test))
```

```
0.21074452488157325
```

```
In [70]: print(en.score(x_train,y_train))
```

```
0.20880071012353074
```

Logistic Regression

```
In [71]: from sklearn.linear_model import LogisticRegression
```

```
In [72]: feature_matrix=df2.iloc[:,0:5]  
target_vector=df2.iloc[:,-1]
```

```
In [73]: from sklearn.preprocessing import StandardScaler
```

```
In [74]: fs=StandardScaler().fit_transform(feature_matrix)
```

```
In [75]: logr=LogisticRegression()  
logr.fit(fs,target_vector)
```

```
C:\Users\Admin\anaconda3\lib\site-packages\sklearn\linear_model\_logistic.py:  
762: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression (https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)
n_iter_i = _check_optimize_result(

```
Out[75]: LogisticRegression()
```

```
In [76]: df2.shape
```

```
Out[76]: (209448, 15)
```

```
In [77]: observation=[[1,2,3,4,5]]  
predication = logr.predict(observation)
```

```
In [78]: print(predication)
```

```
[28079099]
```

```
In [79]: logr.classes_
```

```
Out[79]: array([28079003, 28079004, 28079008, 28079011, 28079016, 28079017,  
                28079018, 28079024, 28079026, 28079027, 28079036, 28079038,  
                28079039, 28079040, 28079047, 28079048, 28079049, 28079050,  
                28079054, 28079055, 28079056, 28079057, 28079058, 28079059,  
                28079060, 28079099], dtype=int64)
```



```
In [80]: from sklearn.model_selection import train_test_split  
x_train,x_test,y_train,y_test = train_test_split(feature_matrix,target_vector,t
```

```
In [81]: print(logr.score(x_test,y_test))
```

0.06501153815548659

```
In [82]: print(logr.score(x_train,y_train))
```

0.06485782297613445

Conclusion ¶

Ridge Regression is bestfit model

The Score x_test,y_test is 0.0.258645747289049

```
In [ ]:
```

```
In [ ]:
```