

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [41]: df=pd.read_csv(r"C:\Users\Admin\Downloads\csvs_per_year\csvs_per_year\madrid_2002.csv")
df
```

Out[41]:

| | date | BEN | CO | EBE | MX | NMHC | NO_2 | NOx | OXY | O_3 | PM10 | PXY | SO_2 |
|--------|---------------------|------|------|------|-------|------|------------|------------|------|-------|-----------|------|-----------|
| 0 | 2002-04-01 01:00:00 | NaN | 1.39 | NaN | NaN | NaN | 145.100006 | 352.100006 | NaN | 6.54 | 41.990002 | NaN | 21.320000 |
| 1 | 2002-04-01 01:00:00 | 1.93 | 0.71 | 2.33 | 6.20 | 0.15 | 98.150002 | 153.399994 | 2.67 | 6.85 | 20.980000 | 2.53 | 11.660000 |
| 2 | 2002-04-01 01:00:00 | NaN | 0.80 | NaN | NaN | NaN | 103.699997 | 134.000000 | NaN | 13.01 | 28.440001 | NaN | 13.670000 |
| 3 | 2002-04-01 01:00:00 | NaN | 1.61 | NaN | NaN | NaN | 97.599998 | 268.000000 | NaN | 5.12 | 42.180000 | NaN | 16.990000 |
| 4 | 2002-04-01 01:00:00 | NaN | 1.90 | NaN | NaN | NaN | 92.089996 | 237.199997 | NaN | 7.28 | 76.330002 | NaN | 15.260000 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 217291 | 2002-11-01 00:00:00 | 4.16 | 1.14 | NaN | NaN | NaN | 81.080002 | 265.700012 | NaN | 7.21 | 36.750000 | NaN | 13.210000 |
| 217292 | 2002-11-01 00:00:00 | 3.67 | 1.73 | 2.89 | NaN | 0.38 | 113.900002 | 373.100006 | NaN | 5.66 | 63.389999 | NaN | 15.640000 |
| 217293 | 2002-11-01 00:00:00 | 1.37 | 0.58 | 1.17 | 2.37 | 0.15 | 65.389999 | 107.699997 | 1.30 | 9.11 | 9.640000 | 0.94 | 5.620000 |
| 217294 | 2002-11-01 00:00:00 | 4.51 | 0.91 | 4.83 | 10.99 | NaN | 149.800003 | 202.199997 | 1.00 | 5.75 | NaN | 5.52 | 24.219999 |
| 217295 | 2002-11-01 00:00:00 | 3.11 | 1.17 | 3.00 | 7.77 | 0.26 | 80.110001 | 180.300003 | 2.25 | 7.38 | 29.240000 | 3.35 | 12.910000 |

217296 rows × 16 columns

```
In [42]: df1 = df.fillna(0)
df1
```

Out[42]:

| | date | BEN | CO | EBE | MXV | NMHC | NO_2 | NOx | OXY | O_3 | PM10 | PXY | SO_2 |
|--------|---------------------|------|------|------|-------|------|------------|------------|------|-------|-----------|------|-----------|
| 0 | 2002-04-01 01:00:00 | 0.00 | 1.39 | 0.00 | 0.00 | 0.00 | 145.100006 | 352.100006 | 0.00 | 6.54 | 41.990002 | 0.00 | 21.320000 |
| 1 | 2002-04-01 01:00:00 | 1.93 | 0.71 | 2.33 | 6.20 | 0.15 | 98.150002 | 153.399994 | 2.67 | 6.85 | 20.980000 | 2.53 | 11.660000 |
| 2 | 2002-04-01 01:00:00 | 0.00 | 0.80 | 0.00 | 0.00 | 0.00 | 103.699997 | 134.000000 | 0.00 | 13.01 | 28.440001 | 0.00 | 13.670000 |
| 3 | 2002-04-01 01:00:00 | 0.00 | 1.61 | 0.00 | 0.00 | 0.00 | 97.599998 | 268.000000 | 0.00 | 5.12 | 42.180000 | 0.00 | 16.990000 |
| 4 | 2002-04-01 01:00:00 | 0.00 | 1.90 | 0.00 | 0.00 | 0.00 | 92.089996 | 237.199997 | 0.00 | 7.28 | 76.330002 | 0.00 | 15.260000 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 217291 | 2002-11-01 00:00:00 | 4.16 | 1.14 | 0.00 | 0.00 | 0.00 | 81.080002 | 265.700012 | 0.00 | 7.21 | 36.750000 | 0.00 | 13.210000 |
| 217292 | 2002-11-01 00:00:00 | 3.67 | 1.73 | 2.89 | 0.00 | 0.38 | 113.900002 | 373.100006 | 0.00 | 5.66 | 63.389999 | 0.00 | 15.640000 |
| 217293 | 2002-11-01 00:00:00 | 1.37 | 0.58 | 1.17 | 2.37 | 0.15 | 65.389999 | 107.699997 | 1.30 | 9.11 | 9.640000 | 0.94 | 5.620000 |
| 217294 | 2002-11-01 00:00:00 | 4.51 | 0.91 | 4.83 | 10.99 | 0.00 | 149.800003 | 202.199997 | 1.00 | 5.75 | 0.000000 | 5.52 | 24.219999 |
| 217295 | 2002-11-01 00:00:00 | 3.11 | 1.17 | 3.00 | 7.77 | 0.26 | 80.110001 | 180.300003 | 2.25 | 7.38 | 29.240000 | 3.35 | 12.910000 |

217296 rows × 16 columns

In [43]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 217296 entries, 0 to 217295
Data columns (total 16 columns):
#   Column      Non-Null Count  Dtype
---  -
0   date        217296 non-null  object
1   BEN         66747 non-null   float64
2   CO          216637 non-null  float64
3   EBE         58547 non-null   float64
4   MXY         41255 non-null   float64
5   NMHC        87045 non-null   float64
6   NO_2        216439 non-null  float64
7   NOx         216439 non-null  float64
8   OXY         41314 non-null   float64
9   O_3         216726 non-null  float64
10  PM10        209113 non-null  float64
11  PXY         41256 non-null   float64
12  SO_2        216507 non-null  float64
13  TCH         87115 non-null   float64
14  TOL         66619 non-null   float64
15  station     217296 non-null  int64
dtypes: float64(14), int64(1), object(1)
memory usage: 26.5+ MB
```

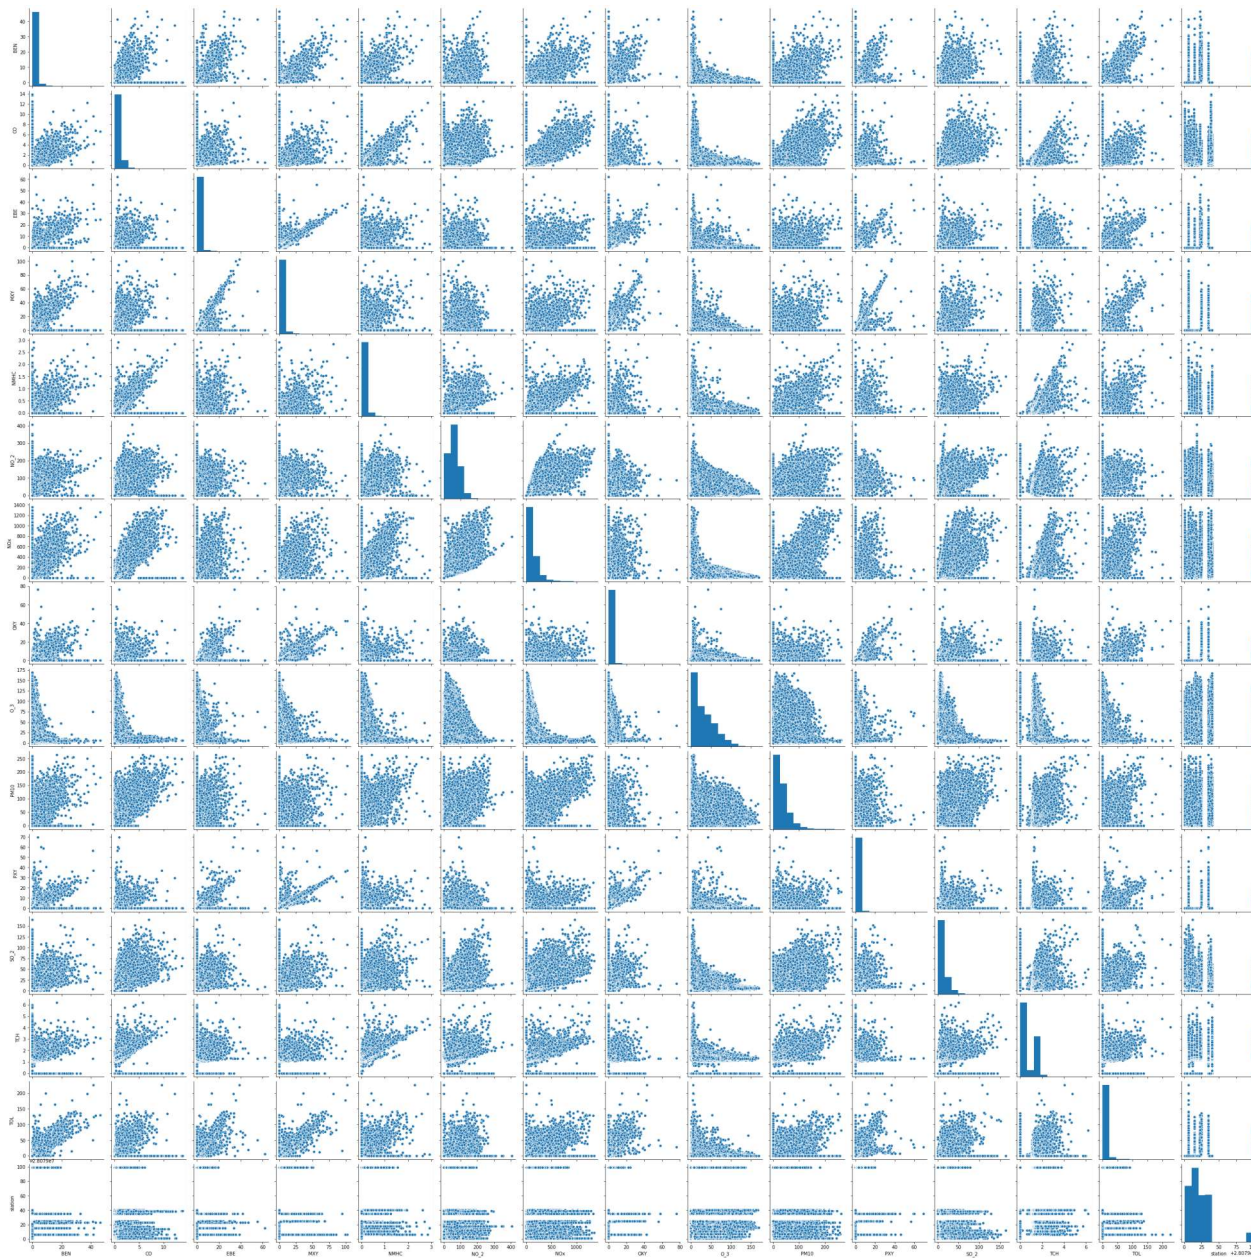
In [44]: `df.columns`

Out[44]: Index(['date', 'BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3', 'PM10', 'PXY', 'SO_2', 'TCH', 'TOL', 'station'], dtype='object')

In [45]: `df2=df1[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3', 'PM10', 'PXY', 'SO_2', 'TCH', 'TOL', 'station']]`


```
In [46]: sns.pairplot(df2)
```

```
Out[46]: <seaborn.axisgrid.PairGrid at 0x23316d258e0>
```




```
In [47]: sns.distplot(df2['O_3'])
```

```
Out[47]: <matplotlib.axes._subplots.AxesSubplot at 0x2333c5d7880>
```



```
In [48]: sns.heatmap(df2.corr())
```

```
Out[48]: <matplotlib.axes._subplots.AxesSubplot at 0x23356ce0280>
```



Linear Regression

```
In [49]: x = df2[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY',
                'PM10', 'PXY', 'SO_2', 'TCH', 'TOL']]
y = df2['O_3']
```

```
In [50]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.3)
```

```
In [51]: from sklearn.linear_model import LinearRegression

lr=LinearRegression()
lr.fit(x_train,y_train)
```

```
Out[51]: LinearRegression()
```

```
In [52]: print(lr.intercept_)
```

57.469790661936074


```
In [53]: coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-effecient'])  
coeff
```

Out[53]:

| | Co-effecient |
|------|--------------|
| BEN | 0.522743 |
| CO | -2.175383 |
| EBE | 0.113789 |
| MXY | 0.070354 |
| NMHC | 14.353484 |
| NO_2 | -0.273541 |
| NOx | -0.079958 |
| OXY | -0.333165 |
| PM10 | 0.211886 |
| PXY | 0.910906 |
| SO_2 | -0.216371 |
| TCH | -2.408368 |
| TOL | -0.243377 |

```
In [54]: prediction=lr.predict(x_test)  
plt.scatter(y_test,prediction)
```

Out[54]: <matplotlib.collections.PathCollection at 0x2331998d100>



```
In [55]: print(lr.score(x_test,y_test))
```

0.3635724046526313

```
In [56]: lr.score(x_train,y_train)
```

Out[56]: 0.36284560662269716

Ridge Lasso

```
In [57]: from sklearn.linear_model import Ridge,Lasso
```

```
In [58]: rr=Ridge(alpha=10)
rr.fit(x_train,y_train)
rr.score(x_test,y_test)
```

```
Out[58]: 0.3635716817834509
```

```
In [59]: predict2=(rr.predict(x_test))
```

```
In [60]: la=Lasso(alpha=10)
la.fit(x_train,y_train)
```

```
Out[60]: Lasso(alpha=10)
```

```
In [61]: la.score(x_test,y_test)
```

```
Out[61]: 0.35790445753202804
```

Elastic Net regression

```
In [62]: from sklearn.linear_model import ElasticNet
en=ElasticNet()
en.fit(x_train,y_train)
```

```
Out[62]: ElasticNet()
```

```
In [63]: print(en.coef_)
```

```
[ 0.         -0.         0.         0.05185668  0.         -0.26186518
 -0.08971544  0.         0.20528086  0.         -0.21073604 -0.
 -0.00487909]
```

```
In [64]: print(en.intercept_)
```

```
56.02023772231269
```

```
In [65]: print(en.score(x_test,y_test))
```

```
0.36070101767965235
```

```
In [66]: print(en.score(x_train,y_train))
```

```
0.3598254732927437
```

Logistic Regression

```
In [67]: from sklearn.linear_model import LogisticRegression
```

```
In [68]: feature_matrix=df2.iloc[:,0:5]
target_vector=df2.iloc[:,-1]
```

```
In [69]: from sklearn.preprocessing import StandardScaler
```



```
In [70]: fs=StandardScaler().fit_transform(feature_matrix)
```

```
In [71]: logr=LogisticRegression()  
logr.fit(fs,target_vector)
```

C:\Users\Admin\anaconda3\lib\site-packages\sklearn\linear_model_logistic.py:762: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression (https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)
n_iter_i = _check_optimize_result(

```
Out[71]: LogisticRegression()
```

```
In [72]: df2.shape
```

```
Out[72]: (217296, 15)
```

```
In [73]: observation=[[1,2,3,4,5]]  
predication = logr.predict(observation)
```

```
In [74]: print(predication)  
  
[28079099]
```

```
In [75]: logr.classes_
```

```
Out[75]: array([28079001, 28079003, 28079004, 28079006, 28079007, 28079008,  
                28079009, 28079011, 28079012, 28079014, 28079015, 28079016,  
                28079017, 28079018, 28079019, 28079021, 28079022, 28079023,  
                28079024, 28079025, 28079035, 28079036, 28079038, 28079039,  
                28079040, 28079099], dtype=int64)
```

```
In [76]: from sklearn.model_selection import train_test_split  
  
x_train,x_test,y_train,y_test = train_test_split(feature_matrix,target_vector,test_size=0.30)
```

```
In [77]: print(logr.score(x_test,y_test))  
  
0.08981576646366718
```

```
In [78]: print(logr.score(x_train,y_train))  
  
0.0886744199806715
```

Conclusion

Linear Regression is bestfit model

Linear Regression is bestfit model for dataset madrid_2001. The Score x_test,y_test is 0.3635724046526313 and x_train,y_train score is 0.36284560662269716

In []:

In []: