

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

In [117]: df=pd.read_csv(r"C:\Users\Admin\Downloads\csvs_per_year\csvs_per_year\madrid_2004.csv")
df
```

Out[117]:

	date	BEN	CO	EBE	MXV	NMHC	NO_2	NOx	OXY	O_3	PM10	PM25	
0	2004-08-01 01:00:00	NaN	0.66	NaN	NaN	NaN	89.550003	118.900002	NaN	40.020000	39.990002	25.860001	↑
1	2004-08-01 01:00:00	2.66	0.54	2.99	6.08	0.18	51.799999	53.860001	3.28	51.689999	22.950001	NaN	↕
2	2004-08-01 01:00:00	NaN	1.02	NaN	NaN	NaN	93.389999	138.600006	NaN	20.860001	49.480000	NaN	↑
3	2004-08-01 01:00:00	NaN	0.53	NaN	NaN	NaN	87.290001	105.000000	NaN	36.730000	31.070000	NaN	↑
4	2004-08-01 01:00:00	NaN	0.17	NaN	NaN	NaN	34.910000	35.349998	NaN	86.269997	54.080002	NaN	↑
...	...	...	...	...	...	...	...	...	...	...	...	...	
245491	2004-06-01 00:00:00	0.75	0.21	0.85	1.55	0.07	59.580002	64.389999	0.66	33.029999	30.900000	14.860000	↕
245492	2004-06-01 00:00:00	2.49	0.75	2.44	4.57	NaN	97.139999	146.899994	2.34	7.740000	37.689999	NaN	↕
245493	2004-06-01 00:00:00	NaN	NaN	NaN	NaN	0.13	102.699997	132.600006	NaN	17.809999	22.840000	12.040000	↑
245494	2004-06-01 00:00:00	NaN	NaN	NaN	NaN	0.09	82.599998	102.599998	NaN	NaN	45.630001	NaN	↑
245495	2004-06-01 00:00:00	3.01	0.67	2.78	5.12	0.20	92.550003	141.000000	2.60	11.460000	24.389999	17.959999	↕

245496 rows × 17 columns

```
In [118]: df1 = df.fillna(0)
df1
```

Out[118]:

	date	BEN	CO	EBE	MXV	NMHC	NO_2	NOx	OXY	O_3	PM10	PM25	F
0	2004-08-01 01:00:00	0.00	0.66	0.00	0.00	0.00	89.550003	118.900002	0.00	40.020000	39.990002	25.860001	(
1	2004-08-01 01:00:00	2.66	0.54	2.99	6.08	0.18	51.799999	53.860001	3.28	51.689999	22.950001	0.000000	3
2	2004-08-01 01:00:00	0.00	1.02	0.00	0.00	0.00	93.389999	138.600006	0.00	20.860001	49.480000	0.000000	(
3	2004-08-01 01:00:00	0.00	0.53	0.00	0.00	0.00	87.290001	105.000000	0.00	36.730000	31.070000	0.000000	(
4	2004-08-01 01:00:00	0.00	0.17	0.00	0.00	0.00	34.910000	35.349998	0.00	86.269997	54.080002	0.000000	(
...	...	...	...	...	...	...	...	...	...	...	...	...	...
245491	2004-06-01 00:00:00	0.75	0.21	0.85	1.55	0.07	59.580002	64.389999	0.66	33.029999	30.900000	14.860000	(
245492	2004-06-01 00:00:00	2.49	0.75	2.44	4.57	0.00	97.139999	146.899994	2.34	7.740000	37.689999	0.000000	2
245493	2004-06-01 00:00:00	0.00	0.00	0.00	0.00	0.13	102.699997	132.600006	0.00	17.809999	22.840000	12.040000	(
245494	2004-06-01 00:00:00	0.00	0.00	0.00	0.00	0.09	82.599998	102.599998	0.00	0.000000	45.630001	0.000000	(
245495	2004-06-01 00:00:00	3.01	0.67	2.78	5.12	0.20	92.550003	141.000000	2.60	11.460000	24.389999	17.959999	2

245496 rows × 17 columns

In [119]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 245496 entries, 0 to 245495
Data columns (total 17 columns):
#   Column      Non-Null Count  Dtype
---  -
0   date        245496 non-null  object
1   BEN         65158 non-null   float64
2   CO          226043 non-null  float64
3   EBE         56781 non-null   float64
4   MXY         39867 non-null   float64
5   NMHC        107630 non-null  float64
6   NO_2        243280 non-null  float64
7   NOx         243283 non-null  float64
8   OXY         39882 non-null   float64
9   O_3         233811 non-null  float64
10  PM10        234655 non-null  float64
11  PM25        58145 non-null   float64
12  PXY         39891 non-null   float64
13  SO_2        243402 non-null  float64
14  TCH         107650 non-null  float64
15  TOL         64914 non-null   float64
16  station     245496 non-null  int64
dtypes: float64(15), int64(1), object(1)
memory usage: 31.8+ MB
```

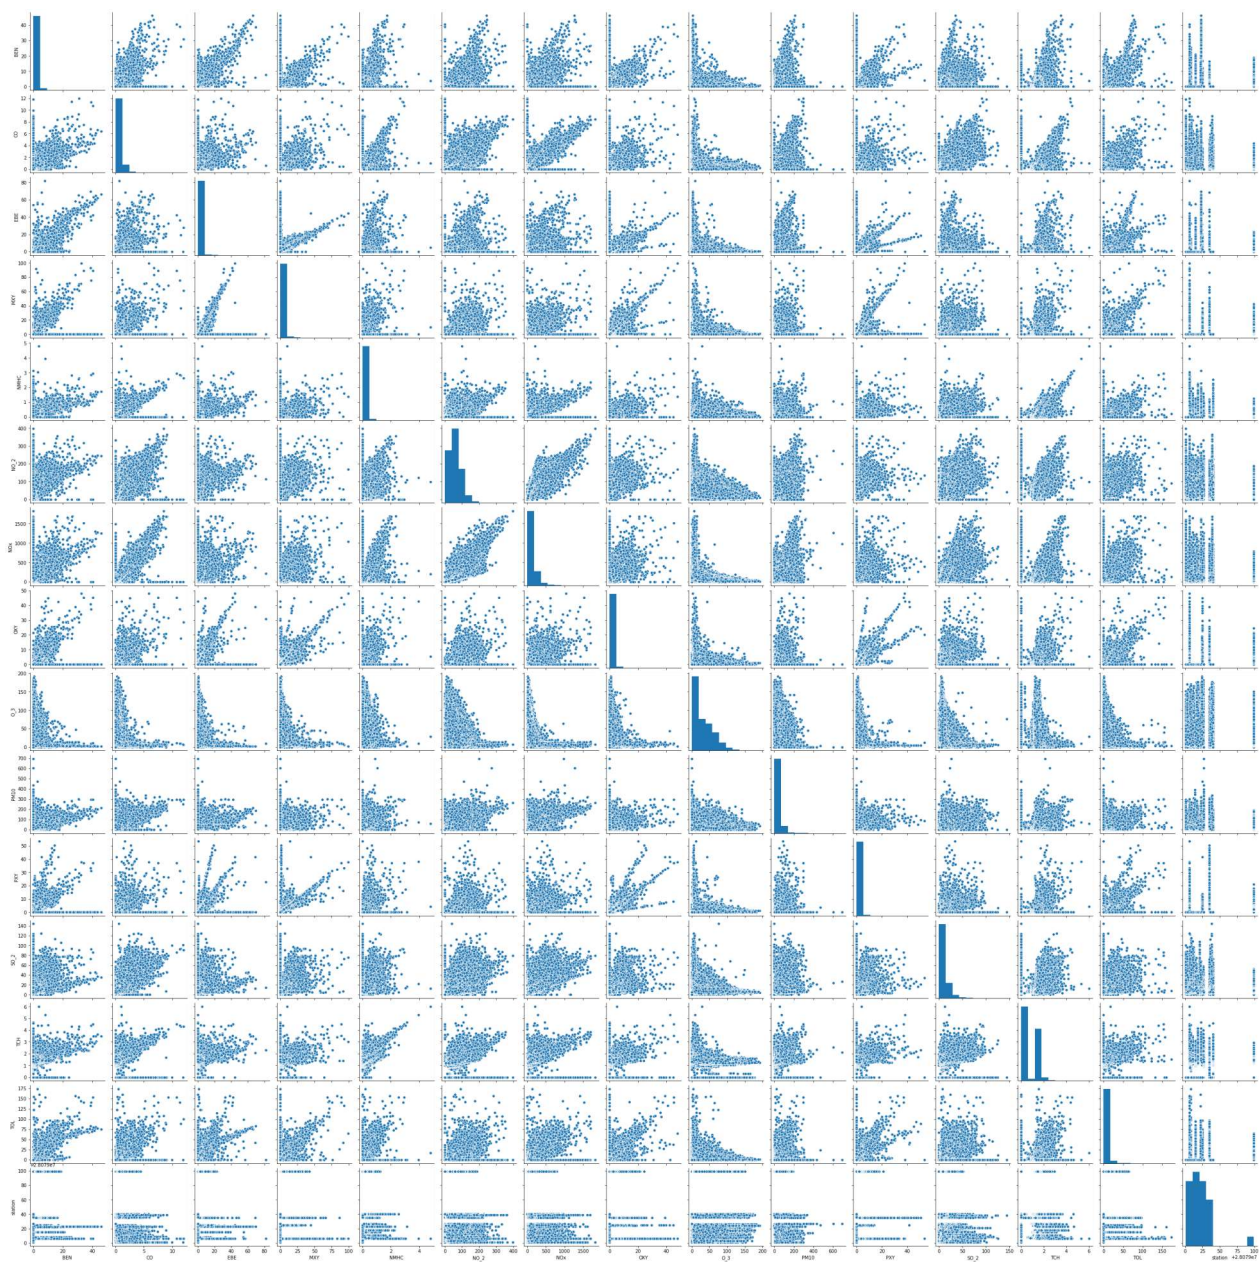
In [120]: `df.columns`

Out[120]: Index(['date', 'BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO\_2', 'NOx', 'OXY', 'O\_3', 'PM10', 'PM25', 'PXY', 'SO\_2', 'TCH', 'TOL', 'station'], dtype='object')

In [121]: `df2=df1[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3', 'PM10', 'PXY', 'SO_2', 'TCH', 'TOL', 'station']]`


```
In [122]: sns.pairplot(df2)
```

Out[122]: <seaborn.axisgrid.PairGrid at 0x233803edf40>




```
In [123]: sns.distplot(df2['O_3'])
```

```
Out[123]: <matplotlib.axes._subplots.AxesSubplot at 0x233cb0e97f0>
```



```
In [124]: sns.heatmap(df2.corr())
```

```
Out[124]: <matplotlib.axes._subplots.AxesSubplot at 0x233cb0da760>
```



## Linear Regression

```
In [125]: x = df2[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY',
                  'PM10', 'PXY', 'SO_2', 'TCH', 'TOL']]
y = df2['O_3']
```

```
In [126]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.3)
```

```
In [127]: from sklearn.linear_model import LinearRegression

lr=LinearRegression()
lr.fit(x_train,y_train)
```

```
Out[127]: LinearRegression()
```

```
In [128]: print(lr.intercept_)
```

```
52.575372173875785
```


```
In [129]: coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-effecient'])
coeff
```

```
Out[129]:
```

	Co-effecient
<b>BEN</b>	1.445308
<b>CO</b>	4.346790
<b>EBE</b>	0.357843
<b>MXV</b>	-0.205967
<b>NMHC</b>	9.029511
<b>NO_2</b>	-0.214476
<b>NOx</b>	-0.092395
<b>OXY</b>	0.500563
<b>PM10</b>	0.080688
<b>PXY</b>	0.537244
<b>SO_2</b>	-0.216725
<b>TCH</b>	-0.353907
<b>TOL</b>	-0.502155

```
In [130]: prediction=lr.predict(x_test)
plt.scatter(y_test,prediction)
```

```
Out[130]: <matplotlib.collections.PathCollection at 0x23343f02a30>
```



```
In [131]: print(lr.score(x_test,y_test))
```

```
0.27848502134774233
```

```
In [132]: lr.score(x_train,y_train)
```

```
Out[132]: 0.27536834626506734
```

## Ridge Lasso

```
In [133]: from sklearn.linear_model import Ridge,Lasso
```

```
In [134]: rr=Ridge(alpha=10)
rr.fit(x_train,y_train)
rr.score(x_test,y_test)
```

```
Out[134]: 0.27848660273191406
```

```
In [135]: predict2=(rr.predict(x_test))
```

```
In [136]: la=Lasso(alpha=10)
la.fit(x_train,y_train)
```

```
Out[136]: Lasso(alpha=10)
```

```
In [137]: la.score(x_test,y_test)
```

```
Out[137]: 0.2707189798655314
```

## Elastic Net regression

```
In [138]: from sklearn.linear_model import ElasticNet
en=ElasticNet()
en.fit(x_train,y_train)
```

```
Out[138]: ElasticNet()
```

```
In [139]: print(en.coef_)
```

```
[ 0.3556417  0.          0.3480369  0.09368703  0.          -0.22289699
 -0.07147375  0.06541692  0.07677801  0.00629613 -0.13864582  0.
 -0.18489377]
```

```
In [140]: print(en.intercept_)
```

```
52.7536621756676
```

```
In [141]: print(en.score(x_test,y_test))
```

```
0.2746437307982825
```

```
In [142]: print(en.score(x_train,y_train))
```

```
0.2702125917708419
```

## Logistic Regression

```
In [143]: from sklearn.linear_model import LogisticRegression
```

```
In [144]: feature_matrix=df2.iloc[:,0:5]
target_vector=df2.iloc[:,-1]
```

```
In [145]: from sklearn.preprocessing import StandardScaler
```

```
In [146]: fs=StandardScaler().fit_transform(feature_matrix)
```

```
In [147]: logr=LogisticRegression()  
logr.fit(fs,target_vector)
```

C:\Users\Admin\anaconda3\lib\site-packages\sklearn\linear\_model\\_logistic.py:762: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:  
<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)  
Please also refer to the documentation for alternative solver options:  
[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression) ([https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression))  
n\_iter\_i = \_check\_optimize\_result(

```
Out[147]: LogisticRegression()
```

```
In [148]: df2.shape
```

```
Out[148]: (245496, 15)
```

```
In [149]: observation=[[1,2,3,4,5]]  
predication = logr.predict(observation)
```

```
In [150]: print(predication)  
  
[28079035]
```

```
In [151]: logr.classes_
```

```
Out[151]: array([28079001, 28079003, 28079004, 28079006, 28079007, 28079008,  
                28079009, 28079011, 28079012, 28079014, 28079015, 28079016,  
                28079017, 28079018, 28079019, 28079021, 28079022, 28079023,  
                28079024, 28079025, 28079026, 28079027, 28079035, 28079036,  
                28079038, 28079039, 28079040, 28079099], dtype=int64)
```

```
In [152]: from sklearn.model_selection import train_test_split  
  
x_train,x_test,y_train,y_test = train_test_split(feature_matrix,target_vector,test_size=0.30
```

```
In [153]: print(logr.score(x_test,y_test))  
  
0.08650490841695067
```

```
In [154]: print(logr.score(x_train,y_train))  
  
0.0847905404225852
```

## Conclusion ¶

Linear Regression is bestfit model

The Score x\_test,y\_test is 0.27848502134774233 and x\_train,y\_train score is 0.27536834626506734



In [ ]:

In [ ]: