

```
In [6]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [7]: df=pd.read_csv(r"C:\Users\Admin\Downloads\csvs_per_year\csvs_per_year\madrid_2008.csv")
df
```

Out[7]:

	date	BEN	CO	EBE	MXV	NMHC	NO_2	NOx	OXY	O_3	PM
0	2009-10-01 01:00:00	NaN	0.27	NaN	NaN	NaN	39.889999	48.150002	NaN	50.680000	18.2600
1	2009-10-01 01:00:00	NaN	0.22	NaN	NaN	NaN	21.230000	24.260000	NaN	55.880001	10.5800
2	2009-10-01 01:00:00	NaN	0.18	NaN	NaN	NaN	31.230000	34.880001	NaN	49.060001	25.1900
3	2009-10-01 01:00:00	0.95	0.33	1.43	2.68	0.25	55.180000	81.360001	1.57	36.669998	26.5300
4	2009-10-01 01:00:00	NaN	0.41	NaN	NaN	0.12	61.349998	76.260002	NaN	38.090000	23.7600
...
215683	2009-06-01 00:00:00	0.50	0.22	0.39	0.75	0.09	22.000000	24.510000	1.00	82.239998	10.8300
215684	2009-06-01 00:00:00	NaN	0.31	NaN	NaN	NaN	76.110001	101.099998	NaN	41.220001	9.9200
215685	2009-06-01 00:00:00	0.13	NaN	0.86	NaN	0.23	81.050003	99.849998	NaN	24.830000	12.4600
215686	2009-06-01 00:00:00	0.21	NaN	2.96	NaN	0.10	72.419998	82.959999	NaN	NaN	13.0300
215687	2009-06-01 00:00:00	0.37	0.32	0.99	1.36	0.14	54.290001	64.480003	1.06	56.919998	15.3600

215688 rows × 12 columns



```
In [8]: df1 = df.fillna(0)
df1
```

Out[8]:

	date	BEN	CO	EBE	MXY	NMHC	NO_2	NOx	OXY	O_3	PM
0	2009-10-01 01:00:00	0.00	0.27	0.00	0.00	0.00	39.889999	48.150002	0.00	50.680000	18.2600
1	2009-10-01 01:00:00	0.00	0.22	0.00	0.00	0.00	21.230000	24.260000	0.00	55.880001	10.5800
2	2009-10-01 01:00:00	0.00	0.18	0.00	0.00	0.00	31.230000	34.880001	0.00	49.060001	25.1900
3	2009-10-01 01:00:00	0.95	0.33	1.43	2.68	0.25	55.180000	81.360001	1.57	36.669998	26.5300
4	2009-10-01 01:00:00	0.00	0.41	0.00	0.00	0.12	61.349998	76.260002	0.00	38.090000	23.7600
...
215683	2009-06-01 00:00:00	0.50	0.22	0.39	0.75	0.09	22.000000	24.510000	1.00	82.239998	10.8300
215684	2009-06-01 00:00:00	0.00	0.31	0.00	0.00	0.00	76.110001	101.099998	0.00	41.220001	9.9200
215685	2009-06-01 00:00:00	0.13	0.00	0.86	0.00	0.23	81.050003	99.849998	0.00	24.830000	12.4600
215686	2009-06-01 00:00:00	0.21	0.00	2.96	0.00	0.10	72.419998	82.959999	0.00	0.000000	13.0300
215687	2009-06-01 00:00:00	0.37	0.32	0.99	1.36	0.14	54.290001	64.480003	1.06	56.919998	15.3600

215688 rows × 17 columns

```
In [9]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 215688 entries, 0 to 215687
Data columns (total 17 columns):
#   Column      Non-Null Count  Dtype
---  -
0   date        215688 non-null  object
1   BEN         60082 non-null   float64
2   CO          190801 non-null  float64
3   EBE         60081 non-null   float64
4   MXY         24846 non-null   float64
5   NMHC        74748 non-null   float64
6   NO_2        214562 non-null  float64
7   NOx         214565 non-null  float64
8   OXY         24854 non-null   float64
9   O_3         204482 non-null  float64
10  PM10        196331 non-null  float64
11  PM25        55822 non-null   float64
12  PXY         24854 non-null   float64
13  SO_2        212671 non-null  float64
14  TCH         75213 non-null   float64
15  TOL         59920 non-null   float64
16  station     215688 non-null  int64
dtypes: float64(15), int64(1), object(1)
memory usage: 28.0+ MB
```

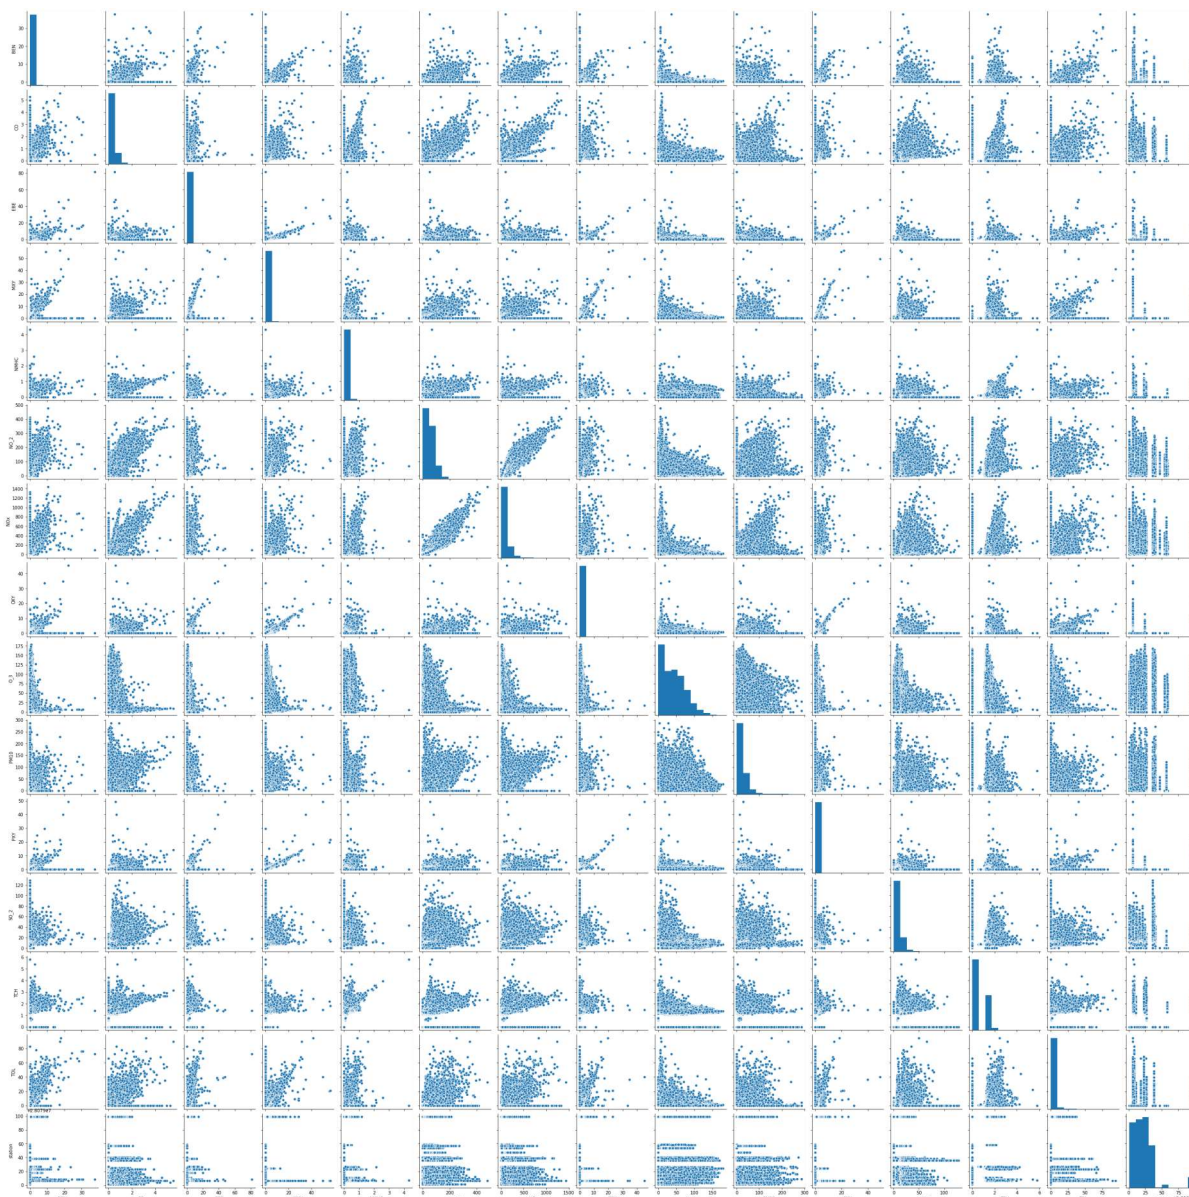
```
In [10]: df.columns
```

```
Out[10]: Index(['date', 'BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
               'PM10', 'PM25', 'PXY', 'SO_2', 'TCH', 'TOL', 'station'],
              dtype='object')
```

```
In [11]: df2=df1[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
                 'PM10', 'PXY', 'SO_2', 'TCH', 'TOL', 'station']]
```


```
In [12]: sns.pairplot(df2)
```

```
Out[12]: <seaborn.axisgrid.PairGrid at 0x1ac70ed5f40>
```




```
In [13]: sns.distplot(df2['O_3'])
```

```
Out[13]: <matplotlib.axes._subplots.AxesSubplot at 0x1ac2e33e7c0>
```



```
In [14]: sns.heatmap(df2.corr())
```

```
Out[14]: <matplotlib.axes._subplots.AxesSubplot at 0x1ac2e3ef160>
```



Linear Regression

```
In [15]: x = df2[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY',
                'PM10', 'PXY', 'SO_2', 'TCH', 'TOL']]
y = df2['O_3']
```

```
In [16]: from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.3)
```

```
In [17]: from sklearn.linear_model import LinearRegression
```

```
lr=LinearRegression()  
lr.fit(x_train,y_train)
```

```
Out[17]: LinearRegression()
```

```
In [18]: print(lr.intercept_)
```

```
66.95558974958072
```


```
In [19]: coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-effecient'])  
coeff
```

```
Out[19]:
```

	Co-effecient
BEN	3.544157
CO	17.404977
EBE	-7.545598
MXY	-2.135929
NMHC	29.425730
NO_2	-0.326915
NOx	-0.113924
OXY	10.360508
PM10	0.171713
PXY	-1.657101
SO_2	-0.583736
TCH	-4.477473
TOL	0.351467

```
In [20]: prediction=lr.predict(x_test)  
plt.scatter(y_test,prediction)
```

```
Out[20]: <matplotlib.collections.PathCollection at 0x1ac2e702910>
```



```
In [21]: print(lr.score(x_test,y_test))
```

```
0.3436569025600422
```

```
In [22]: lr.score(x_train,y_train)
```

```
Out[22]: 0.3327286867692786
```

Ridge Lasso

```
In [23]: from sklearn.linear_model import Ridge,Lasso
```

```
In [24]: rr=Ridge(alpha=10)
rr.fit(x_train,y_train)
rr.score(x_test,y_test)
```

```
Out[24]: 0.3436540606390921
```

```
In [25]: predict2=(rr.predict(x_test))
```

```
In [26]: la=Lasso(alpha=10)
la.fit(x_train,y_train)
```

```
Out[26]: Lasso(alpha=10)
```

```
In [27]: la.score(x_test,y_test)
```

```
Out[27]: 0.2891893530653368
```

Elastic Net regression

```
In [28]: from sklearn.linear_model import ElasticNet
en=ElasticNet()
en.fit(x_train,y_train)
```

```
Out[28]: ElasticNet()
```

```
In [29]: print(en.coef_)
```

```
[ 0.          0.3250155 -2.31668926  0.96889706  0.          -0.33518481
 -0.06882161  0.6987213   0.18557355  0.26107453 -0.3754829  -0.18163418
  0.05852214]
```

```
In [30]: print(en.intercept_)
```

```
66.55187199777671
```

```
In [31]: print(en.score(x_test,y_test))
```

```
0.3093740774829815
```

```
In [32]: print(en.score(x_train,y_train))
```

```
0.3011510507398939
```

Logistic Regression

```
In [33]: from sklearn.linear_model import LogisticRegression
```

```
In [34]: feature_matrix=df2.iloc[:,0:5]  
target_vector=df2.iloc[:,-1]
```

```
In [35]: from sklearn.preprocessing import StandardScaler
```

```
In [36]: fs=StandardScaler().fit_transform(feature_matrix)
```

```
In [37]: logr=LogisticRegression()  
logr.fit(fs,target_vector)
```

```
C:\Users\Admin\anaconda3\lib\site-packages\sklearn\linear_model\_logistic.py:  
762: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression (https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)
n_iter_i = _check_optimize_result(

```
Out[37]: LogisticRegression()
```

```
In [38]: df2.shape
```

```
Out[38]: (215688, 15)
```

```
In [39]: observation=[[1,2,3,4,5]]  
predication = logr.predict(observation)
```

```
In [40]: print(predication)
```

```
[28079099]
```

```
In [41]: logr.classes_
```

```
Out[41]: array([28079001, 28079003, 28079004, 28079006, 28079007, 28079008,  
                28079009, 28079011, 28079012, 28079014, 28079016, 28079017,  
                28079018, 28079019, 28079021, 28079022, 28079023, 28079024,  
                28079025, 28079026, 28079027, 28079036, 28079038, 28079039,  
                28079040, 28079047, 28079054, 28079057, 28079058, 28079059,  
                28079099], dtype=int64)
```



```
In [42]: from sklearn.model_selection import train_test_split  
x_train,x_test,y_train,y_test = train_test_split(feature_matrix,target_vector,t
```

```
In [43]: print(logr.score(x_test,y_test))
```

0.07248056624476486

```
In [44]: print(logr.score(x_train,y_train))
```

0.0703797166530888

Conclusion

Ridge Regression is bestfit model

The Score x_test,y_test is 0.3436540606390921

```
In [ ]:
```

```
In [ ]:
```