

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [269]: df=pd.read_csv(r"C:\Users\Admin\Downloads\csvs_per_year\csvs_per_year\madrid_2008\madrid_2008.csv")
df
```

Out[269]:

	date	BEN	CO	EBE	MXV	NMHC	NO_2	NOx	OXY	O_3	PI
0	2008-06-01 01:00:00	NaN	0.47	NaN	NaN	NaN	83.089996	120.699997	NaN	16.990000	16.889
1	2008-06-01 01:00:00	NaN	0.59	NaN	NaN	NaN	94.820000	130.399994	NaN	17.469999	19.040
2	2008-06-01 01:00:00	NaN	0.55	NaN	NaN	NaN	75.919998	104.599998	NaN	13.470000	20.270
3	2008-06-01 01:00:00	NaN	0.36	NaN	NaN	NaN	61.029999	66.559998	NaN	23.110001	10.850
4	2008-06-01 01:00:00	1.68	0.80	1.70	3.01	0.30	105.199997	214.899994	1.61	12.120000	37.160
...	...	...	...	...	...	...	...	...	...	...	...
226387	2008-11-01 00:00:00	0.48	0.30	0.57	1.00	0.31	13.050000	14.160000	0.91	57.400002	5.450
226388	2008-11-01 00:00:00	NaN	0.30	NaN	NaN	NaN	41.880001	48.500000	NaN	35.830002	15.020
226389	2008-11-01 00:00:00	0.25	NaN	0.56	NaN	0.11	83.610001	102.199997	NaN	14.130000	17.540
226390	2008-11-01 00:00:00	0.54	NaN	2.70	NaN	0.18	70.639999	81.860001	NaN	NaN	11.910
226391	2008-11-01 00:00:00	0.75	0.36	1.20	2.75	0.16	58.240002	74.239998	1.64	31.910000	12.690

226392 rows × 17 columns



```
In [270]: df1 = df.fillna(0)
df1
```

Out[270]:

	date	BEN	CO	EBE	MXV	NMHC	NO_2	NOx	OXY	O_3	PI
0	2008-06-01 01:00:00	0.00	0.47	0.00	0.00	0.00	83.089996	120.699997	0.00	16.990000	16.889
1	2008-06-01 01:00:00	0.00	0.59	0.00	0.00	0.00	94.820000	130.399994	0.00	17.469999	19.040
2	2008-06-01 01:00:00	0.00	0.55	0.00	0.00	0.00	75.919998	104.599998	0.00	13.470000	20.270
3	2008-06-01 01:00:00	0.00	0.36	0.00	0.00	0.00	61.029999	66.559998	0.00	23.110001	10.850
4	2008-06-01 01:00:00	1.68	0.80	1.70	3.01	0.30	105.199997	214.899994	1.61	12.120000	37.160
...	...	...	...	...	...	...	...	...	...	...	...
226387	2008-11-01 00:00:00	0.48	0.30	0.57	1.00	0.31	13.050000	14.160000	0.91	57.400002	5.450
226388	2008-11-01 00:00:00	0.00	0.30	0.00	0.00	0.00	41.880001	48.500000	0.00	35.830002	15.020
226389	2008-11-01 00:00:00	0.25	0.00	0.56	0.00	0.11	83.610001	102.199997	0.00	14.130000	17.540
226390	2008-11-01 00:00:00	0.54	0.00	2.70	0.00	0.18	70.639999	81.860001	0.00	0.000000	11.910
226391	2008-11-01 00:00:00	0.75	0.36	1.20	2.75	0.16	58.240002	74.239998	1.64	31.910000	12.690

226392 rows × 17 columns



In [271]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 226392 entries, 0 to 226391
Data columns (total 17 columns):
#   Column      Non-Null Count  Dtype
---  ---
0   date        226392 non-null object
1   BEN         67047 non-null  float64
2   CO          208109 non-null float64
3   EBE         67044 non-null  float64
4   MXY         25867 non-null  float64
5   NMHC        85079 non-null  float64
6   NO_2        225315 non-null float64
7   NOx         225311 non-null float64
8   OXY         25878 non-null  float64
9   O_3         215716 non-null float64
10  PM10        220179 non-null float64
11  PM25        67833 non-null  float64
12  PXY         25877 non-null  float64
13  SO_2        225405 non-null float64
14  TCH         85107 non-null  float64
15  TOL         66940 non-null  float64
16  station     226392 non-null int64
dtypes: float64(15), int64(1), object(1)
memory usage: 29.4+ MB
```

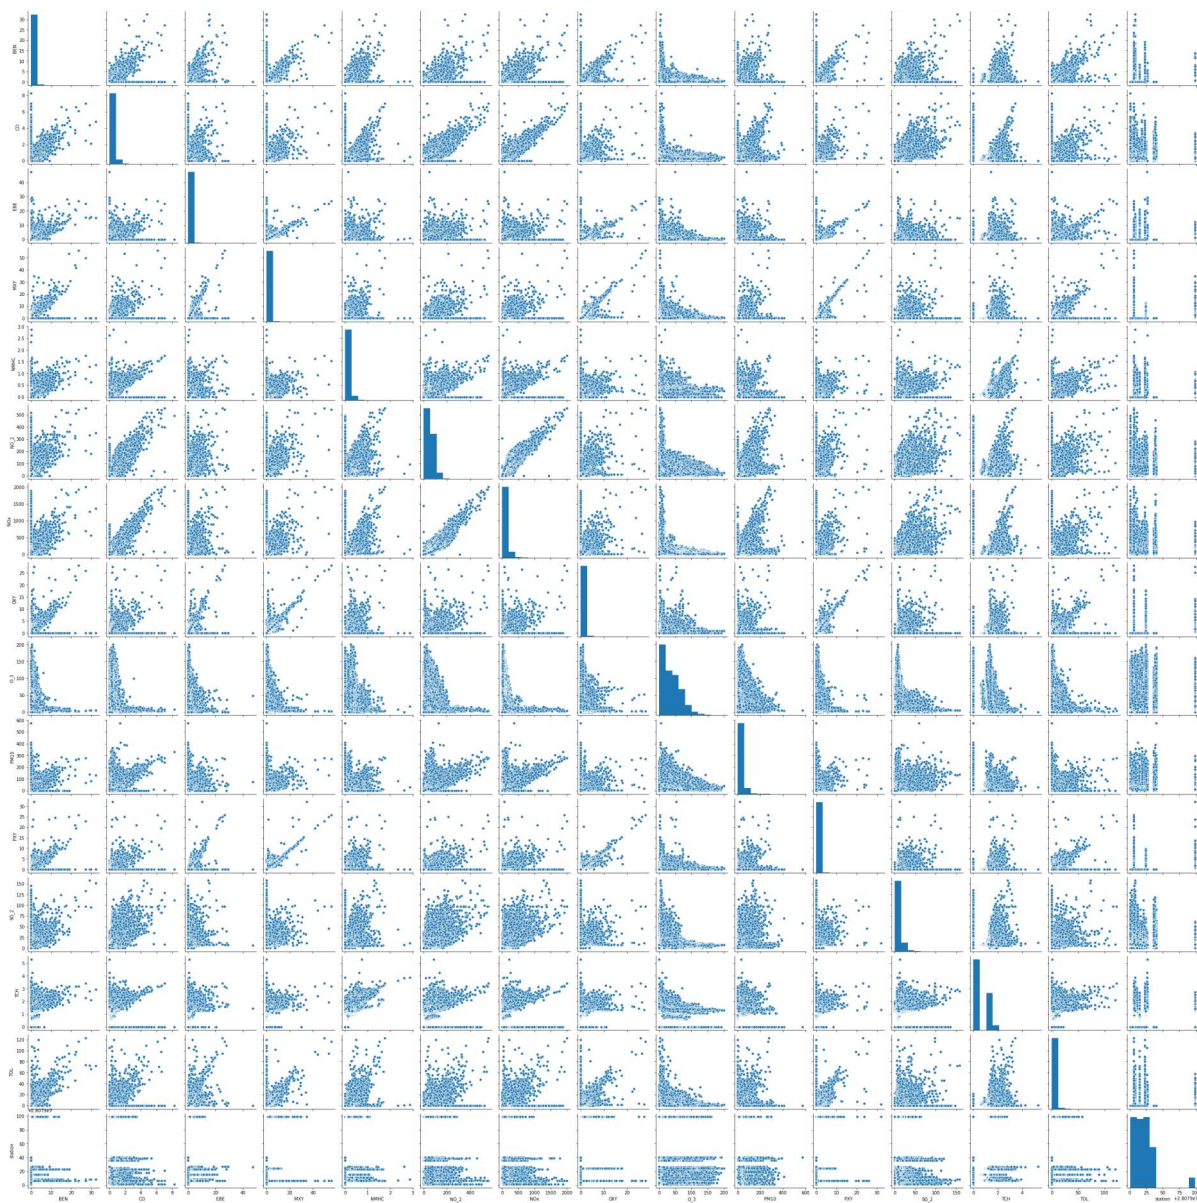
In [272]: df.columns

```
Out[272]: Index(['date', 'BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
                'PM10', 'PM25', 'PXY', 'SO_2', 'TCH', 'TOL', 'station'],
              dtype='object')
```

```
In [273]: df2=df1[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
                  'PM10', 'PXY', 'SO_2', 'TCH', 'TOL', 'station']]
```


```
In [274]: sns.pairplot(df2)
```

```
Out[274]: <seaborn.axisgrid.PairGrid at 0x233e3762460>
```




```
In [275]: sns.distplot(df2['O_3'])
```

```
Out[275]: <matplotlib.axes._subplots.AxesSubplot at 0x2350b2fb4f0>
```



```
In [276]: sns.heatmap(df2.corr())
```

```
Out[276]: <matplotlib.axes._subplots.AxesSubplot at 0x2350dd62910>
```



## Linear Regression

```
In [277]: x = df2[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY',
                  'PM10', 'PXY', 'SO_2', 'TCH', 'TOL']]
y = df2['O_3']
```

```
In [278]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.3)
```

```
In [279]: from sklearn.linear_model import LinearRegression
```

```
lr=LinearRegression()  
lr.fit(x_train,y_train)
```

```
Out[279]: LinearRegression()
```

```
In [280]: print(lr.intercept_)
```

```
59.08351350855042
```


```
In [281]: coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-effecient'])  
coeff
```

```
Out[281]:
```

	Co-effecient
<b>BEN</b>	3.804848
<b>CO</b>	16.892690
<b>EBE</b>	-1.916293
<b>MXY</b>	-2.446715
<b>NMHC</b>	-2.624988
<b>NO_2</b>	-0.282340
<b>NOx</b>	-0.099939
<b>OXY</b>	3.144582
<b>PM10</b>	0.097670
<b>PXY</b>	4.304703
<b>SO_2</b>	-0.512151
<b>TCH</b>	-1.584020
<b>TOL</b>	-0.091192

```
In [282]: prediction=lr.predict(x_test)
plt.scatter(y_test,prediction)
```

Out[282]: <matplotlib.collections.PathCollection at 0x23511da4dc0>



```
In [283]: print(lr.score(x_test,y_test))
```

0.30472025600643393

```
In [284]: lr.score(x_train,y_train)
```

Out[284]: 0.2973802470196698

## Ridge Lasso

```
In [285]: from sklearn.linear_model import Ridge,Lasso
```

```
In [286]: rr=Ridge(alpha=10)
rr.fit(x_train,y_train)
rr.score(x_test,y_test)
```

Out[286]: 0.30472134365526515

```
In [287]: predict2=(rr.predict(x_test))
```

```
In [288]: la=Lasso(alpha=10)
la.fit(x_train,y_train)
```

Out[288]: Lasso(alpha=10)

```
In [289]: la.score(x_test,y_test)
```

Out[289]: 0.2802924270178049

## Elastic Net regression

```
In [290]: from sklearn.linear_model import ElasticNet  
en=ElasticNet()  
en.fit(x_train,y_train)
```

Out[290]: ElasticNet()

```
In [291]: print(en.coef_)
```

```
[ 0.0162072  0.         -0.         0.59858908 -0.         -0.32385527  
 -0.04199357  0.40519487  0.09301885  0.14153153 -0.35292605 -0.59160507  
  0.03162004]
```

```
In [292]: print(en.intercept_)
```

60.34758566862013

```
In [293]: print(en.score(x_test,y_test))
```

0.2885227128727964

```
In [294]: print(en.score(x_train,y_train))
```

0.2806035137518861

## Logistic Regression

```
In [295]: from sklearn.linear_model import LogisticRegression
```

```
In [296]: feature_matrix=df2.iloc[:,0:5]  
target_vector=df2.iloc[:, -1]
```

```
In [297]: from sklearn.preprocessing import StandardScaler
```

```
In [298]: fs=StandardScaler().fit_transform(feature_matrix)
```



```
In [299]: logr=LogisticRegression()  
logr.fit(fs,target_vector)
```

C:\Users\Admin\anaconda3\lib\site-packages\sklearn\linear\_model\\_logistic.py:  
762: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:  
<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)  
Please also refer to the documentation for alternative solver options:  
[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression) ([https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression))  
n\_iter\_i = \_check\_optimize\_result(

```
Out[299]: LogisticRegression()
```

```
In [300]: df2.shape
```

```
Out[300]: (226392, 15)
```

```
In [301]: observation=[[1,2,3,4,5]]  
predication = logr.predict(observation)
```

```
In [302]: print(predication)
```

```
[28079099]
```

```
In [303]: logr.classes_
```

```
Out[303]: array([28079001, 28079003, 28079004, 28079006, 28079007, 28079008,  
                28079009, 28079011, 28079012, 28079014, 28079015, 28079016,  
                28079018, 28079019, 28079021, 28079022, 28079023, 28079024,  
                28079025, 28079026, 28079027, 28079036, 28079038, 28079039,  
                28079040, 28079099], dtype=int64)
```

```
In [304]: from sklearn.model_selection import train_test_split  
  
x_train,x_test,y_train,y_test = train_test_split(feature_matrix,target_vector,t
```

```
In [305]: print(logr.score(x_test,y_test))
```

```
0.10528873052798964
```

```
In [306]: print(logr.score(x_train,y_train))
```

```
0.10626348801696177
```

## Conclusion

Linear Regression is bestfit model

The Score  $x_{\text{test}}, y_{\text{test}}$  is 0.30472025600643393 and  $x_{\text{train}}, y_{\text{train}}$  score is 0.2973802470196698

In [ ]:

In [ ]: